AI DELIVERY RISK ASSISTANT

Grounded delivery risk auditing — not another status summarizer.

Green on the outside. Red on the inside.

Project 3 | Akansha Verma

Why this problem stays invisible today

The “watermelon” status report

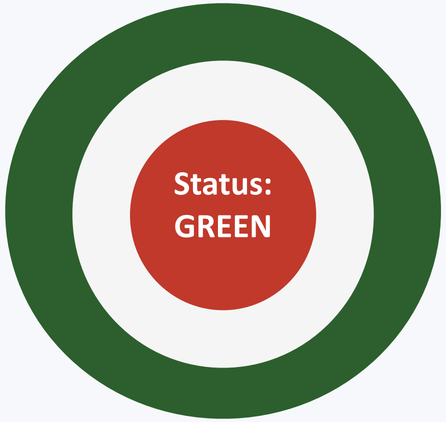

“On track — no blockers to flag.”

- A weekly status report says the project is on track.
- A ticket export shows a critical dependency blocked for 9 days.
- A standup transcript has an engineer saying they don't expect to hit the deadline.

All of that information already exists — nobody has connected it.

Why existing approaches miss it

1. **Jira / Linear dashboards**

   Systems of record. Show what was entered — not whether it's still true, and never cross-reference against a transcript or postmortem.

2. **Generic LLM chat**

   Paste documents, get a summary. No structural guarantee against fabricating a claim or blending unrelated context.

3. **Manual PM synthesis**

   The actual status quo. Reading everything by hand, with no structural check for contradictions.

How: three non-negotiable design principles

1 Citation-or-reject grounding

Every risk traces to a real cited chunk, enforced in code. Unsupported claims are rejected before reporting.

2 Context-consistency guardrail

Two individually-true facts from different projects can't be stitched into one risk. Cross-project contamination is rejected.

3 Deterministic escalation, human approval

SEV-1 findings and status contradictions are never auto-published — a rule-based decision tree routes them to a human via Telegram.

System architecture

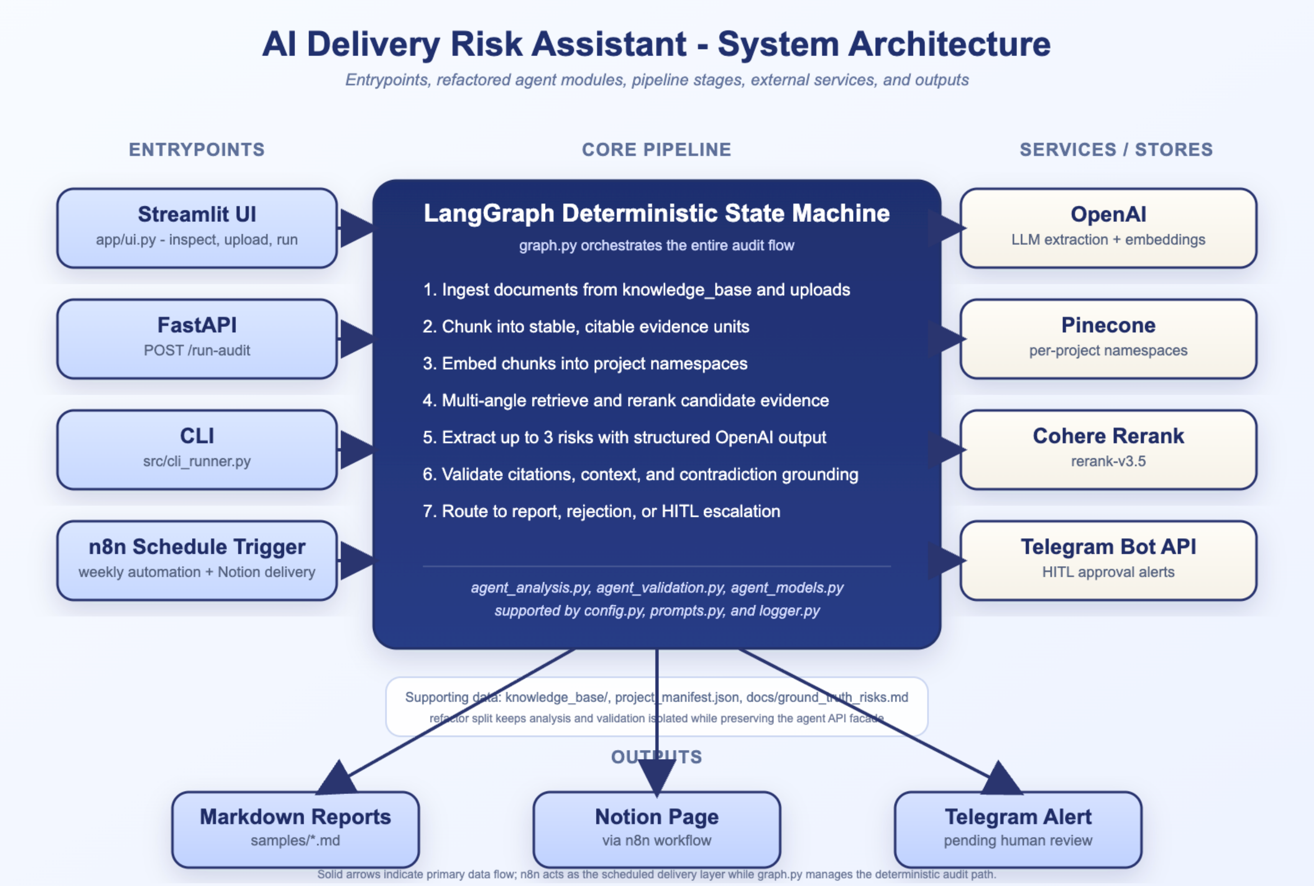

Evaluated against a planted ground-truth answer key

Not just “does it run” — does it surface the right risks, citing the right evidence.

Correctly triggered human escalation

On the single hardest planted case — a contradiction requiring evidence from 4 separate documents.

Under-citation found, root-caused, fixed

Initially cited 2 of 4 sources. A prompt edit alone had no effect — evidence inspection proved the model had all 4 and chose not to cite them. Fixed with a deterministic ticket-ID guardrail; now cites 4, verified against a real re-run.

Zero hallucinated risks

Across both negative controls: a resolved issue and a fabricated budget/vendor scenario — neither appeared.

Where this goes next

Sprint 1 — Pilot Validation

Buyer 1–2 engineering managers

Success metric
Manual synthesis time reduced; ≥1 risk caught early

Sprint 2 — Slack & Jira Native

Buyer Eng managers, tech leads, VP Eng

Success metric
40% weekly active usage; <5 min onboarding

Sprint 3 — Portfolio Dashboard

Buyer CTO, VP Delivery, Eng Ops Director

Success metric
3 paid pilot conversions; 25% less scope creep

Not “can AI summarize project information?”

It's: “How do we safely use AI to improve operational decisions — without treating the AI itself as the source of truth?”

Thank you — questions welcome